

Views and Indexes

Instructor: David Martin

Reference:

*A First Course in Database Systems,
3rd edition, Chapter 2.3 and 8.1-8.4
(except 8.2.3 and 8.4.3)*

SQL Language

- Data Definition Language (DDL)
 - Modify structure of data
 - CREATE, DROP, ALTER
- Data Manipulation Language (DML)
 - Access and modify data
 - SELECT, INSERT, DELETE, UPDATE
- Data Control Language (DCL)
 - Control access to the data (security)
 - GRANT, REVOKE
- Databases also have Utilities, such as Backup/Restore
 - Syntax not specified in the SQL standard

VIEWS: Motivation for Views

- Views help with logical data independence, allowing you to retrieve data as if it matched the description in the view.

```
CREATE VIEW < view-name> AS < view-definition> ;
```

```
CREATE VIEW ParamountMovies AS  
    SELECT title, year  
    FROM Movies  
    WHERE studioName = 'Paramount ' ;
```

- You may now ask queries on ParamountMovies as if it were a table:
 SELECT title FROM ParamountMovies WHERE year=1976;
 – Composition in SQL is powerful: Tables, Queries, Views

More Views

Movies (title , year , length , genre , studioName , producerC #)

MovieExec (name , address , cert # , netWorth)

```
CREATE VIEW MovieProd AS
```

```
    SELECT title, name, genre
```

```
    FROM Movies , MovieExec
```

```
    WHERE producerC # = cert # ;
```

```
SELECT DISTINCT genre
```

```
FROM MovieProd
```

```
WHERE name = 'George Lucas';
```

Renaming Attributes in CREATE VIEW

Movies (title , year , length , genre , studioName , producerC #)

MovieExec (name , address , cert # , netWorth)

```
CREATE VIEW MovieProd(movie_title, prod_name, movie_genre) AS
```

```
    SELECT title, name, genre
```

```
    FROM Movies , MovieExec
```

```
    WHERE producerC# = cert# ;
```

```
SELECT DISTINCT movie_genre
```

```
FROM MovieProd
```

```
WHERE prod_name = 'George Lucas';
```

What is a View?

- A view can include any SQL SELECT statement
 - Including UNION, Aggregates, GROUP BY, HAVING, ORDER BY, etc.
- A view is not stored as a table
 - The tables underlying the view are stored in the database, but only the description of the view is in the database
 - ... although some systems support MATERIALIZED VIEWS
- But view can be used in many (not all) of the same ways as tables
 - Views can be queried
 - Views can be defined on views, as well as on tables!

Queries on Views and Tables

```
CREATE VIEW ParamountMovies AS  
    SELECT title , year  
    FROM Movies  
    WHERE studioName = 'Paramount ' ;
```

```
SELECT DISTINCT starName  
FROM ParamountMovies , StarsIn  
WHERE title = movieTitle AND year = movieYear ;
```

```
CREATE VIEW ParamountStars AS  
    SELECT DISTINCT starName  
    FROM ParamountMovies , StarsIn  
    WHERE title = movieTitle AND year = movieYear ;
```

DROP VIEW

```
CREATE VIEW ParamountMovies AS  
    SELECT title , year  
    FROM Movies  
    WHERE studioName = 'Paramount ' ;
```

```
DROP View ParamountMovies;
```

- What happens if you execute the following?
 - SELECT * FROM ParamountMovies;
 - SELECT * FROM Movies;

View Updates

- Modifications (INSERT/DELETE/UPDATE) are allowed on certain views with simple definitions
 - SELECT query involving a single relation R
 - R not involved in a subquery
 - Attributes of R not appearing in the SELECT clause are assignable with NULL or a default
- Modifications on more complex views are disallowed because:
 - Constraint on underlying table would be violated, or
 - View update is not well-defined
- The SQL rules are complex
- Example for Movies(title, year, length, genre, studioName , producerC#)

```
CREATE VIEW ParamountMovies AS  
    SELECT title , year  
    FROM Movies  
    WHERE studioName = 'Paramount ' ;
```

```
INSERT INTO ParamountMovies VALUES ('StarTrek', 1979);
```

- OK if the other columns of Movies (besides title and year) have defaults or allow nulls

Types of Relations in an RDB

- Three types of relations
 1. Base relations (stored relations, tables)
 - These are tables that contain tuples and can be modified or queried.
 2. Views (derived relations, virtual relations)
 - Views are relations that are defined in terms of other relations but they are not stored. They are constructed only when needed.
 3. Temporary tables
 - These are tables (representing intermediate results) that are constructed by the query execution engine during the processing of a query and discarded when done.

INDEXES: Motivation for Indexes

- Searching an entire table may take a long time:

```
SELECT *
```

```
FROM Movies
```

```
WHERE studioName = 'Disney' AND year = 1990;
```

If there were 100 Million movies, searching them might take a while. An index (e.g., a B-Tree) would allow faster access to matching movies.

If a table is updated, all relevant Indexes on that table are immediately automatically modified within the same transaction.

```
SELECT *  
FROM Movies  
WHERE Title = Monsters, Inc.  
AND Year = 1990
```

Primary Index

MovieIndex

Title	Year	Ptr
Alien	1979	5
Back to the Future	1985	6
Jurassic Park	1998	3
Life Is Beautiful	1997	8
Monsters, Inc.	1990	2
Pretty Woman	1990	1
Princess Mononoke	1997	7
Star Wars IV	1977	4

Primary key

Movies

Title	Year	Length	Genre	Studio	...
Princess Mononoke	1997	134	Fantasy	DENTSU	...
Monsters Inc.	1990	121	Animation	Dreamworks	...
Jurassic Park	1998	145	Adventure	Disney	...
Star Wars IV	1977	121	Sci-fi	LucasFilm	...
Alien	1979	117	Sci-fi	20 th Century Fox	
Back to the Future	1985	116	Sci-fi	Universal	
Pretty Woman	1990	119	Romantic	Disney	...
Life Is Beautiful	1997	116	Comedy	Melampo	

Binary search

Data Structures and Algorithms

- Computer science offers a variety of data structures and algorithms to support indexing
 - E.g., binary search trees, B-trees, B+trees and associated algorithms
 - Note: a great deal of research and progress fundamental computer science has been driven by database requirements

```
SELECT *
FROM Movies
WHERE Year = 1990
AND Studio = Disney
```

Secondary Indexes (1)

YearIndex		<div>Primary key</div> <div>Movies</div>							
Year	Ptr	Title	Year	Length	Genre	Studio	...	Ptr	Studio
1977	4	Princess Mononoke	1997	134	Fantasy	DENTSU	...	5	20 th Century Fox
1979	5	Monsters Inc.	1990	121	Animation	Dreamworks	...	1	DENTSU
1985	6	Jurassic Park	1998	145	Adventure	Disney	...	3	Disney
1990	2	Star Wars IV	1977	121	Sci-fi	LucasFilm	...	7	Disney
1990	1	Alien	1979	117	Sci-fi	20 th Century Fox	...	2	DreamWorks
1997	8	Back to the Future	1985	116	Sci-fi	Universal	...	4	LucasFilm
1997	7	Pretty Woman	1990	119	Romantic	Disney	...	8	Melampo
1998	3	Life Is Beautiful	1997	116	Comedy	Melampo	...	6	Universal

```
SELECT *
FROM Movies
WHERE Year = 1990
AND Studio = Disney
```

Secondary Indexes (2)

YearIndex		<div>Primary key</div> <div>Movies</div>							
Year	Ptr	Title	Year	Length	Genre	Studio	...	Ptr	Studio
1977	4	Princess Mononoke	1997	134	Fantasy	DENTSU	...	5	20 th Century Fox
1979	5	Monsters Inc.	1990	121	Animation	Dreamworks	...	1	DENTSU
1985	6	Jurassic Park	1998	145	Adventure	Disney	...	3	Disney
1990	2	Star Wars IV	1977	121	Sci-fi	LucasFilm	...	7	Disney
1990	1	Alien	1979	117	Sci-fi	20 th Century Fox		2	DreamWorks
1997	8	Back to the Future	1985	116	Sci-fi	Universal		4	LucasFilm
1997	7	Pretty Woman	1990	119	Romantic	Disney	...	8	Melampo
1998	3	Life Is Beautiful	1997	116	Comedy	Melampo		6	Universal

CREATE INDEX

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney' AND year = 1990;
```

How much would each of these indexes help?

```
CREATE INDEX YearIndex ON Movies(year);
```

```
CREATE INDEX StudioIndex ON Movies(studioName);
```

```
CREATE INDEX SYIndex ON Movies(studioName,year);
```

How much would each of the indexes help if the WHERE clause was just
year = 1990?

Disadvantages of Indexes?

- Why not put indexes on every attribute, or even on every combination of attributes that you might query on?
 - Huge number of indexes
 - Space for indexes
 - Cache impact of searching indexes
 - Update time for indexes when table is modified

Selection of Indexes

- Selecting indexes involves evaluation of a trade-off:
 - Great speed up for the affected queries, vs.
 - Significant performance penalties
- Common choices for indexing
 - Primary key
 - Used frequently
 - At most one (data) *page* will be retrieved
 - Attribute that's "almost a key"
 - Relatively few tuples for a given value for that attribute
 - Attribute for which values are frequently specified in queries
 - Attribute for which tuples are "clustered"
 - Common values for the attribute are stored nearby (on a small number of pages)
- Depends mostly on what queries are used heavily

Index Utilization

- SQL statements don't specify use of indexes, so they don't have to be modified when you change what's indexed!
 - Database Optimizer tries to figure out "the best"/"a good" way to execute each SQL query.
 - All the tuples in a Relation can be scanned directly, without using indexes, so indexes aren't necessary ... except for performance.
 - Some systems have ways that you can tell the Optimizer what to do. This has advantages and disadvantages. (What are they?)
- Many SQL systems (including PostgreSQL) have an EXPLAIN PLAN statement, so that you can see what plan the optimizer chooses for a SQL statement.